



Reference alignment of reads

Richard Orton

MRC-University of Glasgow Centre for Virus Research

March 2025

Richard.Orton@glasgow.ac.uk

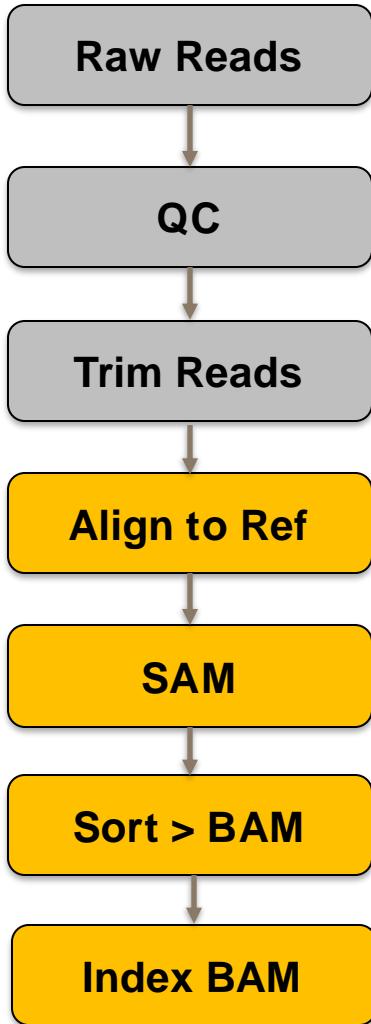


The way and the truth and the life

Here is the bird that never flew
Here is the tree that never grew
Here is the bell that never rang
Here is the fish that never swam



Previously ...



- Previous session we learnt about FASTQ reads and read cleaning/trimming
- Task now is to align these reads to a selected reference sequence

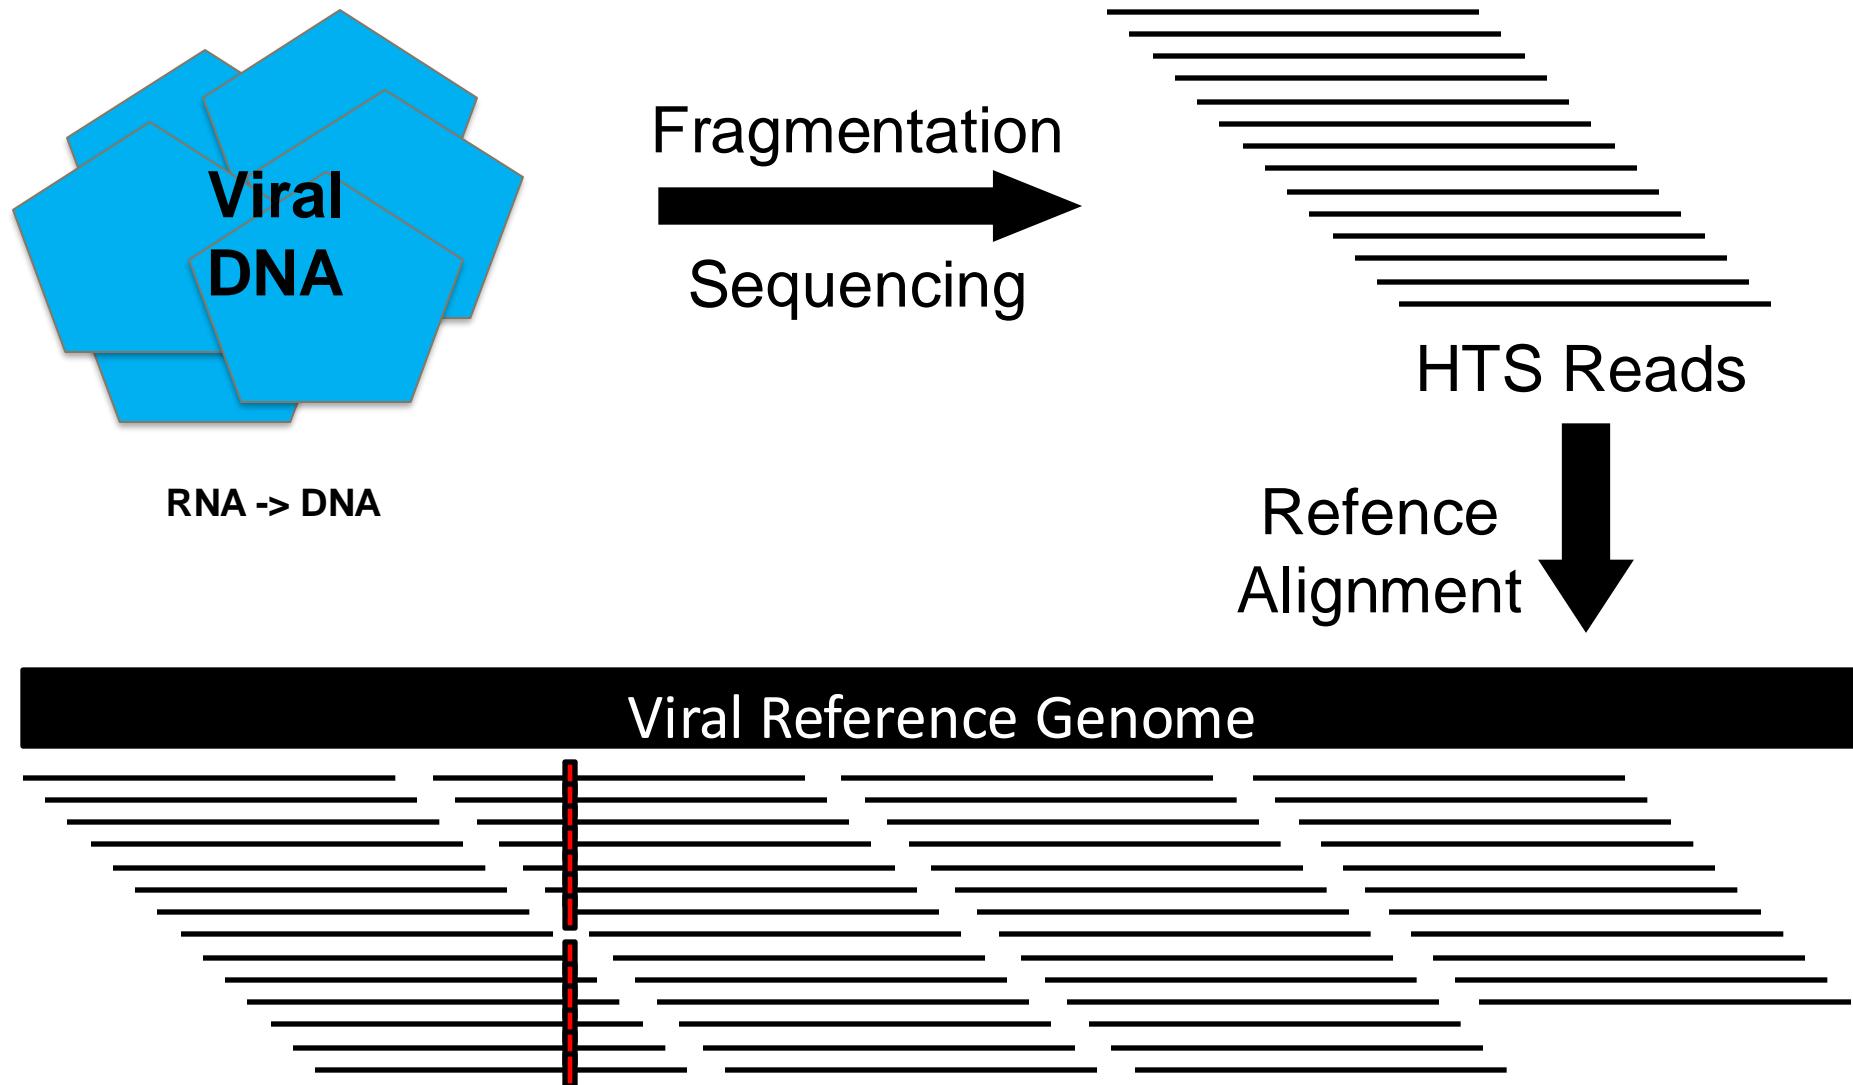
Overview

- **What is reference alignment?**
 - How does it work?
- **What tools can you use?**
- **What do the results look like?**
 - Basic statistics
 - Coverage plots
- **Reference alignment practical**
 - Learn the basic steps of refence alignment, SAM/BAM conversion, calculating basic mapping statistics and coverage plots.

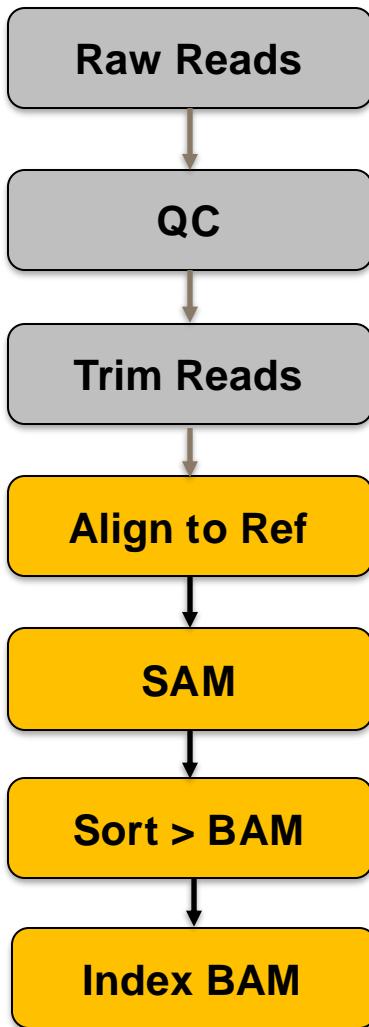
Reference alignment

- **Reference alignment:** want to know the **exact position** on the genome a read originates
 - And the **base-to-base** correspondence (to extract mutations, indels)
- **Reference assembly:** assemble reads back together to form a genome
 - Assemble from scratch – *de novo* assembly – using read overlaps, kmers

Aligning reads to a reference genome



Ref alignment basic steps



```
trim_galore -q 25 --length 50 --paired R1.fq R2.fq
```

```
bwa index ref.fasta  
bwa mem -r ref.fasta R1_qc.fq R2_qc.fq > my.sam
```

```
samtools sort my.sam -o my.bam
```

```
samtools index my.bam
```

Aligning reads to a reference - needs

- **Need Reads**
 - Single or paired, short or long
 - Typically pre-trimmed & filtered
 - But you can use your raw read files
- **Need a Reference**
 - A suitable reference
 - [More on this later]
- Trimmed reads were aligned to the HCV reference genome (GenBank accession NC_038882) with BWA {Li et, 2009}.

Be careful – aligners tend not to complain

- **Sample**
 - Ebola virus sample from a human patient
- **Reads**
 - Reads were adapter trimmed and quality filtered using trim_galore (quality 25, length 50).
- **Reference**
 - Reads were aligned to the HCV reference genome (GenBank accession NC_038882)
- **Result – SAM file of all the reads aligned to the reference**
 - **No errors**
 - **Number of mapped reads (0), coverage statistics (0 cov)**

Unmapped reads

- Reads that could not be aligned to the reference sequence are marked as unmapped
- What are these reads?

Unmapped reads

- Reads that could not be aligned to the reference sequence are marked as unmapped
- What are these reads?
 - Host
 - Bacteria, Parasites,
 - Other viruses
 - Random "low complexity" sequences
- This will be missed as we are “targeting” a specific reference sequence to align against
 - Possible solution: metagenomics

Aligning reads to a references

Ref: ACGGTGACACGTAGCAGTACGC GGTTACACAGA

Read: GTTACAC

Aligning reads to a references

Ref: ACGGTGACACGTAGCAGTACGCAGGTTACACAGA
Read: GTTACAC

↓

↓

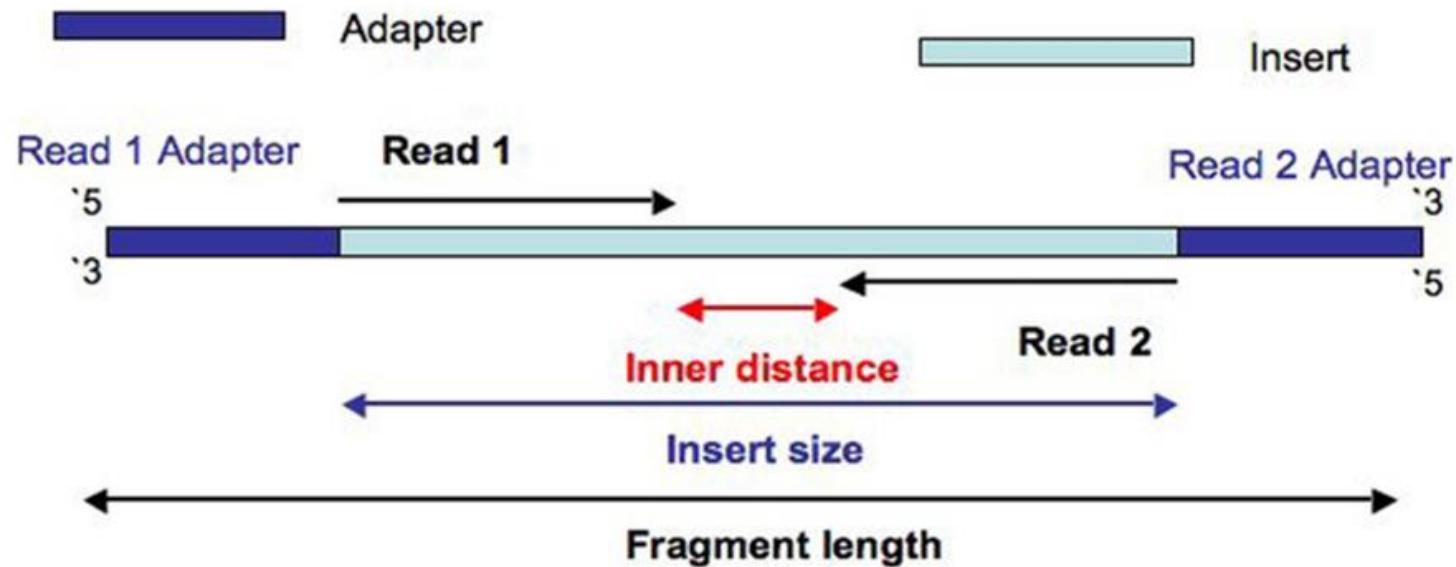
Matches 0/7 7/7

Mismatches 7/7 0/7

Aligners check the reverse complement

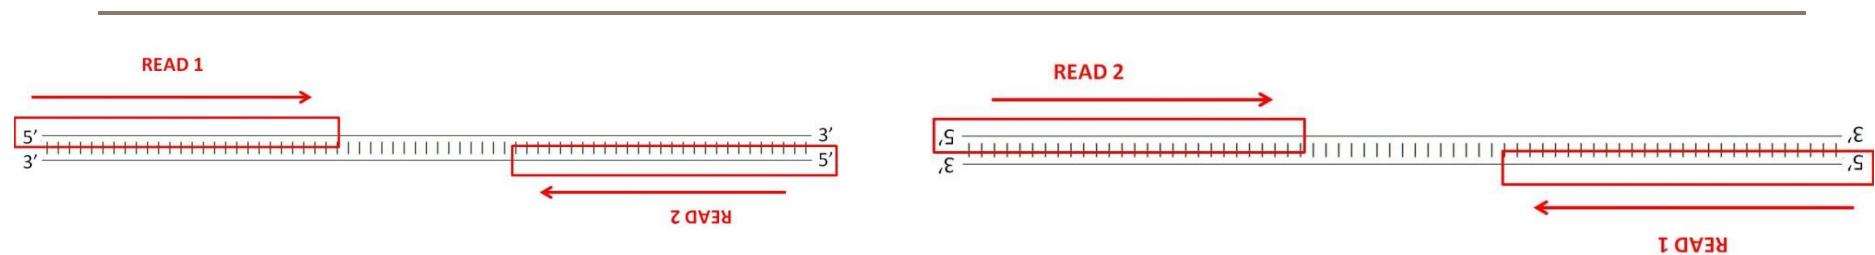
Ref:	ACGGTGACACGTAGCAGTACGCAGGTTACACAGA		
Read:	ACTGCTA	TAGCAGT	
	↓	↓	
			ACTGCTA
			Reverse Complement
Matches	4/7	7/7	TAGCAGT
Mismatches	3/7	0/7	

Paired end ... Insert Size



- Typically, the two reads do not overlap, but they can if the reads are long and fragments are short [redundant data, but can be used to correct errors]
- 500bp fragments + 2 x 300bp reads = 100bp overlap
- Turner 2014, Frontiers in Genetics

Concordance & Discordance – paired reads



Concordant: consistent orientation of read pairs with respect to reference, have insert size within the expected range (depends on library)



Discordant: inconsistent orientation (mixed up reads? reference issues? Abnormal insert size)



Aligning reads to a reference: Mutations and Indels

Ref: ACGGTGACACGTAGCAGTACGC~~GGG~~TTACACAGA
 ACGG~~C~~GA CAGT~~T~~CG AC-~~C~~AGA
 ~~A~~GACGTA GCGGGTT
 TTACACAG
 ~~G~~CGACAC TCGCGGG
 CGG~~C~~GAC AGT~~T~~CGC TACACAT
 ACG-~~A~~GC GGG~~G~~TAC

CIGAR

Concise Idiosyncratic Gapped Alignment Report

1

2

3

Pos: 1234567890123456789012345678901234

Ref: ACGGTGACACGTAGCAGTACGCGGGTTACACAGA

ACGGCGA

CAGTT~~T~~CG

AC-CAGA

~~A~~GACGTA

GC~~GGG~~TT

TTACACAG

~~G~~CGACAC

~~T~~CGCGGG

CGG~~C~~GAC

AGTT~~T~~CGC

TACACAT~~T~~

ACG-AGC

GGG~~G~~TAC

Cov: 12233343333323333334333344433331

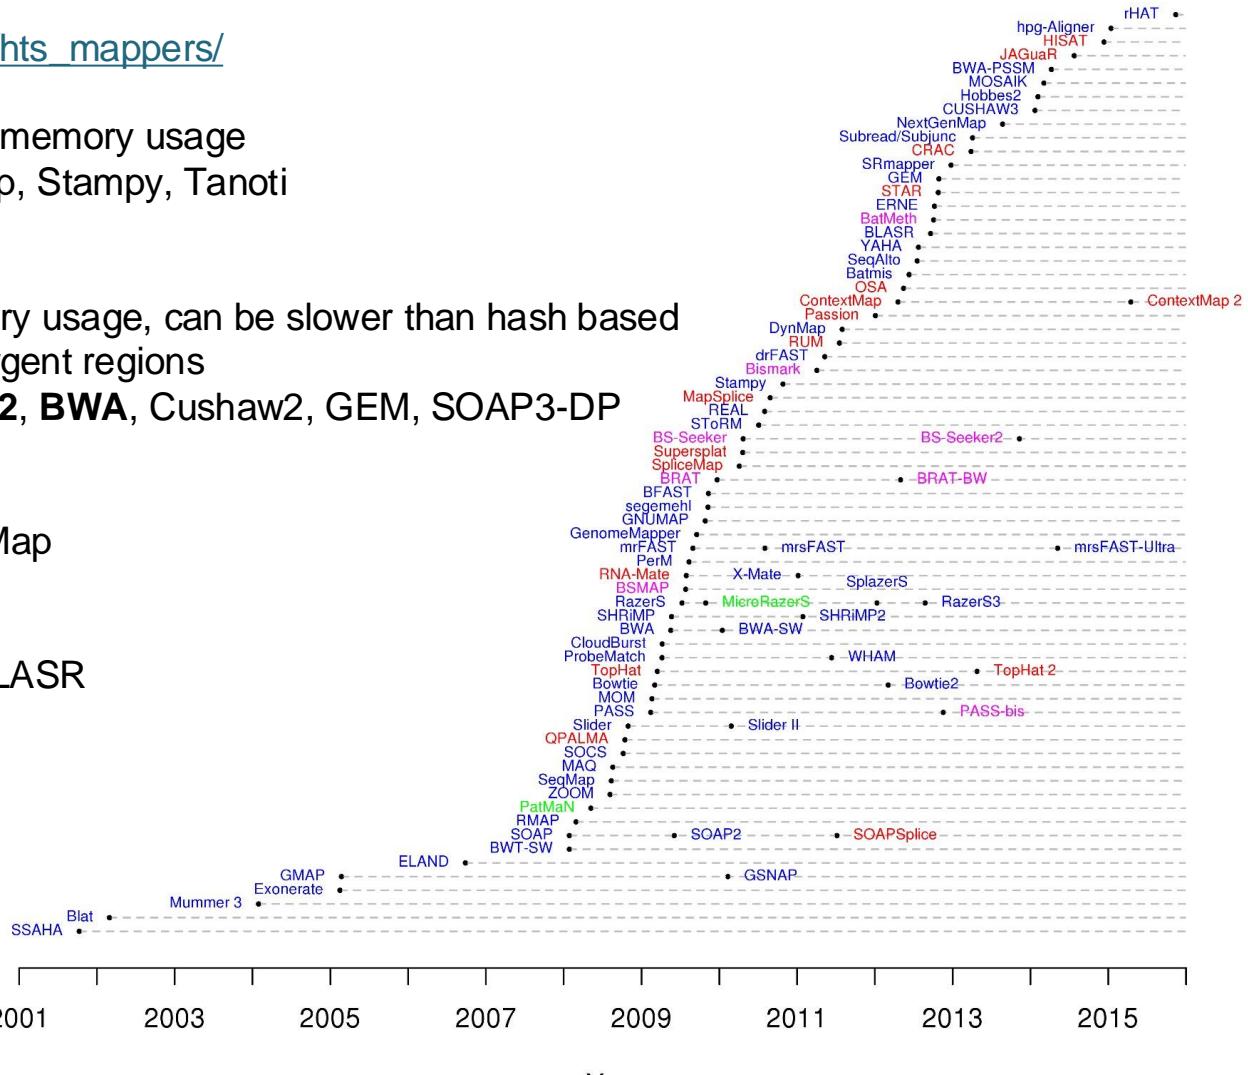
CIGAR

Pos4: 1M1X5M

Pos28: 2M1D4M

Aligners – There are Lots

- https://www.ebi.ac.uk/~nf/hts_mappers/
 - Hash based - faster, high memory usage
 - Mosaik, NextGenMap, Stumpy, Tanoti
 - Burrows-Wheeler based
 - Sensitive, low memory usage, can be slower than hash based
 - Can struggle in divergent regions
 - BarraCUDA, **Bowtie2**, **BWA**, Cushaw2, GEM, SOAP3-DP
 - RNA-Seq Splice aware
 - HiSAT, TopHat, BBMap
 - Long Reads
 - **Minimap2**, LAST, BLASR



Which aligner to use?

Bowtie2

- Sequencing technology – long vs short reads

BWA

- Library/Analysis – e.g. rna-seq

Tanoti

- **Short RNA viral genome - which aligner?**

- In general aligners are quite consistent in terms of consensus sequence & coverage to a good (close) reference

BBMAP

- Differences in aligner can be subtle – so may influence **low frequency** variants

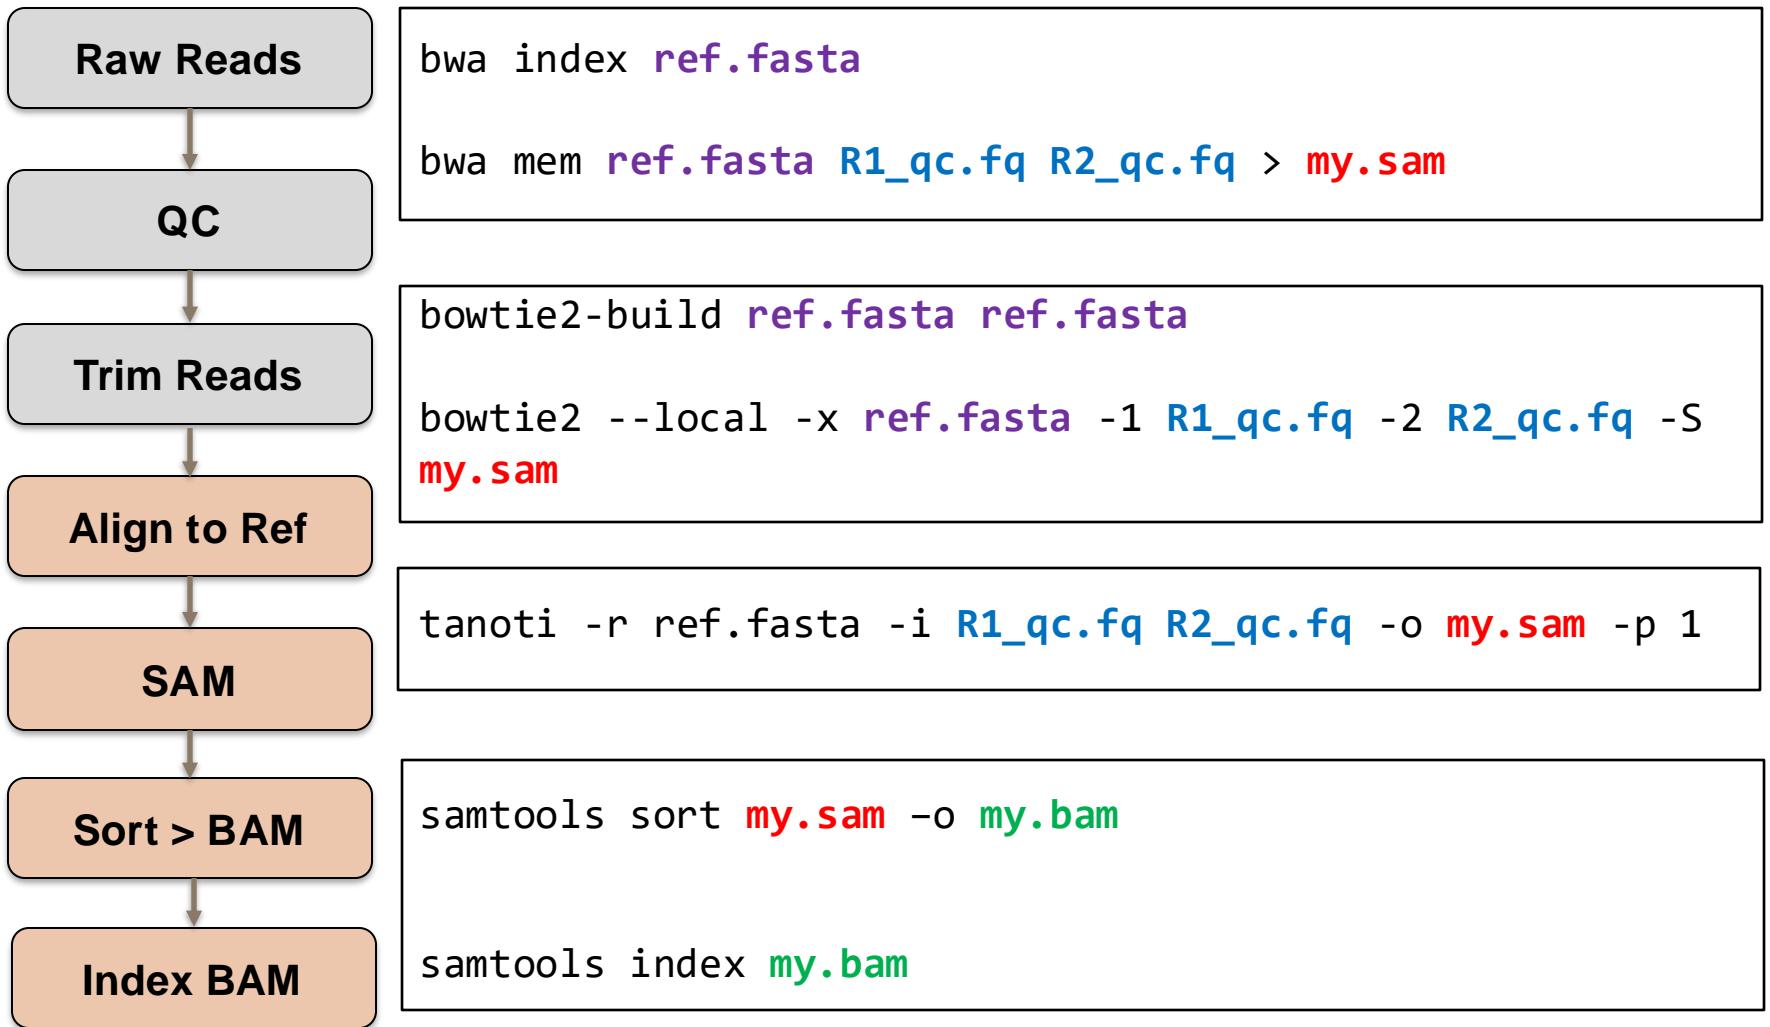
minimap2

- Starting out on a new virus - try a few aligners – not just about most reads aligned – consensus seq and variants

Mosaik

...

Ref alignment commands - different tools



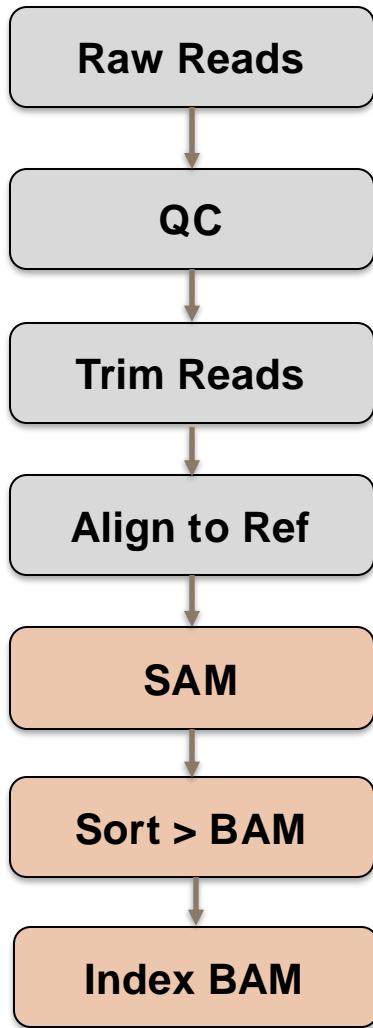
Which reference sequence?

- **Selecting a suitable reference sequence is an important step.**
 - If your reference is too divergent from your sample it can affect read mapping and possibly the consensus sequence
- **First – obviously want to select the right virus!**
 - If you doing a reference assembly – you probably suspect a particular virus is present in your sample
- Second – if a divergent virus e.g. HCV – select the right genotype:
 - Hepatitis C Virus (HCV) - want to select the right genotype – differ by 30–35% at the nucleotide level (subtypes can differ by 15-25% at nucleotide level)
- If unsure what virus is in the sample or suspect it is very divergent
 - **De novo assembly**
 - SHIVER (HIV)
 - Kraken
 - Panel alignment to all genotypes/subtypes – check stats

Multiple Reference sequences

- The reference is in FASTA format: **Need not be one sequence**
 - Segmented virus
 - Influenza: PB1, PB2, PA, NP, HA, NA, M, NS
 - Host
 - Human chromosome 1, 2, 3, 4, 5 etc
 - Panel of viruses
 - HCV 1a, 1b, 1c, 2a, 2b
 - Respiratory viruses
 - Contigs from metagenomics

SAM & BAM files



The result of the alignment step is typically a SAM file

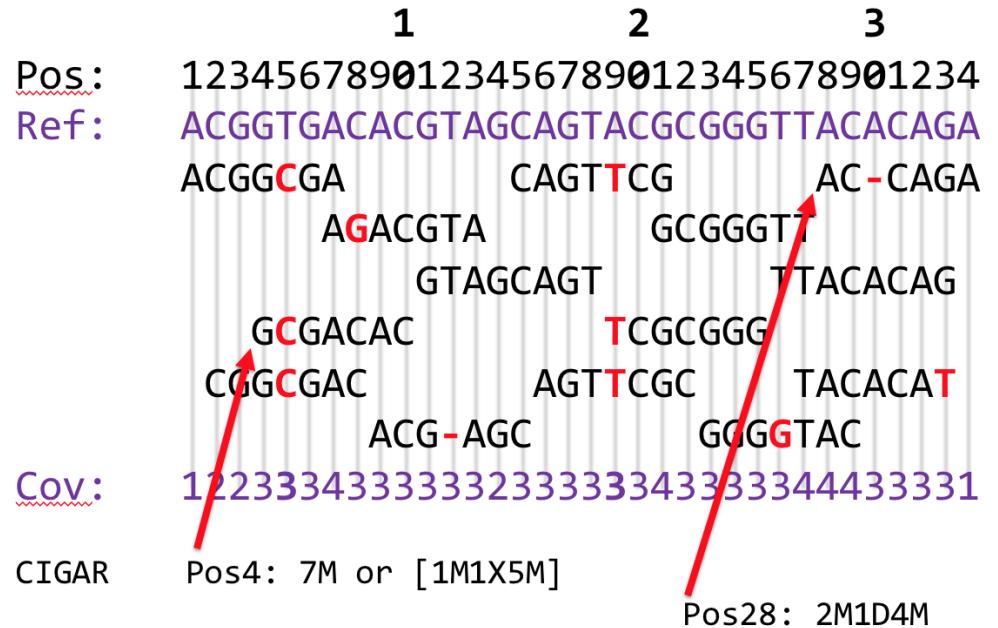
This is then sorted and converted to a BAM file, and indexed

```
samtools sort my.sam -o my.bam
```

```
samtools index my.bam
```

SAM files: Sequence Alignment MAP

- Virtually all aligners output results in **SAM** format
 - Sequence Alignment/Map**
- Each line in the SAM file corresponds to a separate alignment
- Sequence and quality strings of the reads stored in the BAM
 - Can extract reads back out of SAM/BAM
 - But always keeps copies of your raw data



The name of the other read in pair



The position the other read in pair is aligned



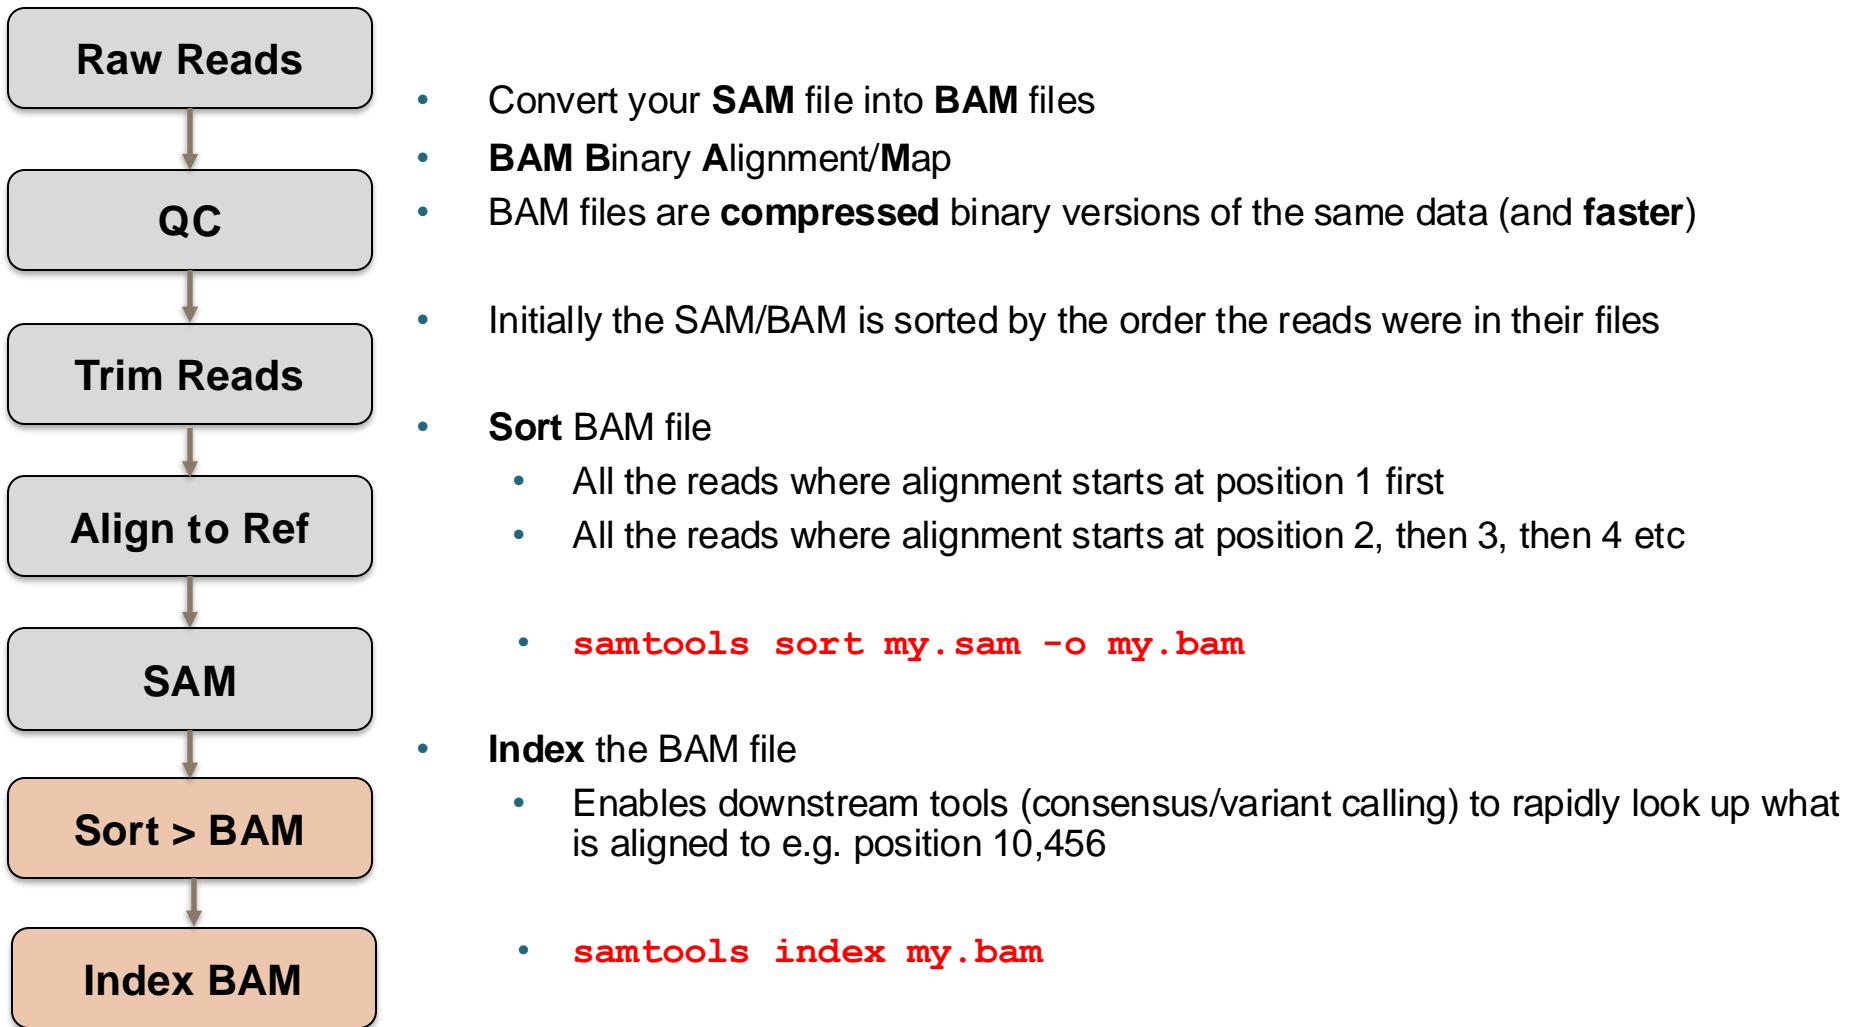
Template Length/Insert Size

QNAME	FLAG	RNAME	POS	MAPQ	CIGAR	RNEXT	PNEXT	TLEN	SEQ	QUALITY
Read3	10	MyRefSeq	28	52	2M1D4M				ACCAGA	IHGFFF
Read8	10	MyRefSeq	4	57	1M1X5M				GCGACAC	IIIIHGGG

Samtools

- **One of the key HTS programs** - provides various utilities for manipulating alignments in the SAM/BAM [and CRAM] formats
 - sorting, merging, indexing and generating alignments in a per-position format.
- Links seamlessly to downstream tools such as VCFTools, BCFTools etc

Converting SAMs to BAMs



SAM Flags – Mapped/Unmapped

- **4 = Read unmapped**
- Can be used to give you the most basic of statistics – how many reads are mapped to the reference and how many are unmapped
- Technically, it is counting how many mapped read alignments are in the SAM file

#	Flag	Description
1	1	Read paired
2	2	Read mapped in proper pair
3	4	Read unmapped
4	8	Mate unmapped
5	16	Read reverse strand
6	32	Mate reverse strand
7	64	First in pair
8	128	Second in pair
9	256	Not primary alignment
10	512	Read fails platform/vendor quality checks
11	1024	Read is PCR or optical duplicate
12	2048	Supplementary alignment

 SAM Flag = 2nd field of SAM file

QNAME	FLAG	RNAME	POS	MAPQ	CIGAR	RNEXT	PNEXT	TLEN	SEQ	QUALITY
Read3	10	MyRefSeq	28	52	2M1D4M				ACCAGA	IHGFFF
Read8	4	*	0	0					GCGACAC	IIHHGGG

SAM Flags – Mapped/Unmapped

- A read can sometimes have multiple alignments
- **256** = not primary = secondary = alternative alignments (equally good or not quite as good)
- **2048** = supplementary alignment = when read is split (spliced) and sections aligned separately

#	Flag	Description
1	1	Read paired
2	2	Read mapped in proper pair
3	4	Read unmapped
4	8	Mate unmapped
5	16	Read reverse strand
6	32	Mate reverse strand
7	64	First in pair
8	128	Second in pair
9	256	Not primary alignment
10	512	Read fails platform/vendor quality checks
11	1024	Read is PCR or optical duplicate
12	2048	Supplementary alignment

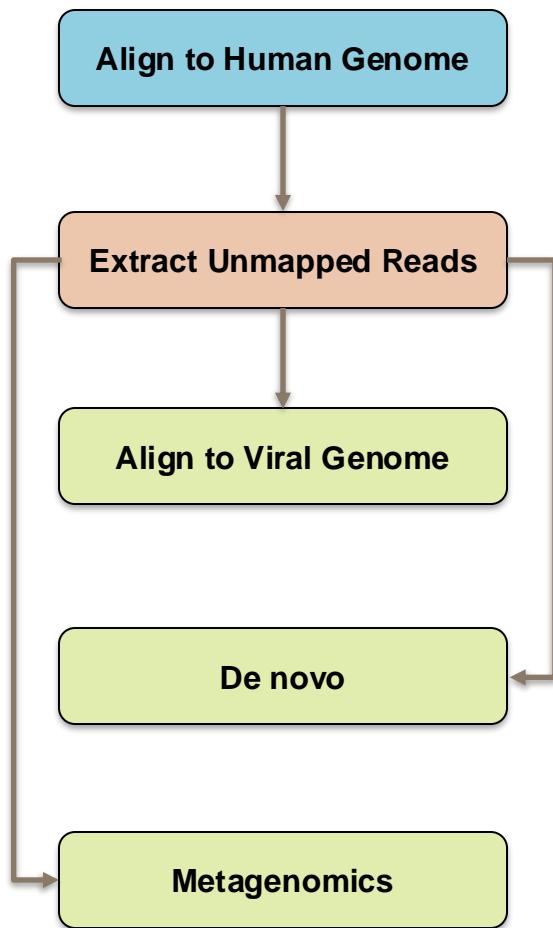
Typically, secondary/supplementary should be few for short RNA virus genome, but if lots it can indicate:

Repeat regions, Large deletions, Poor reference sequence

SAM Flag = 2nd field of SAM file

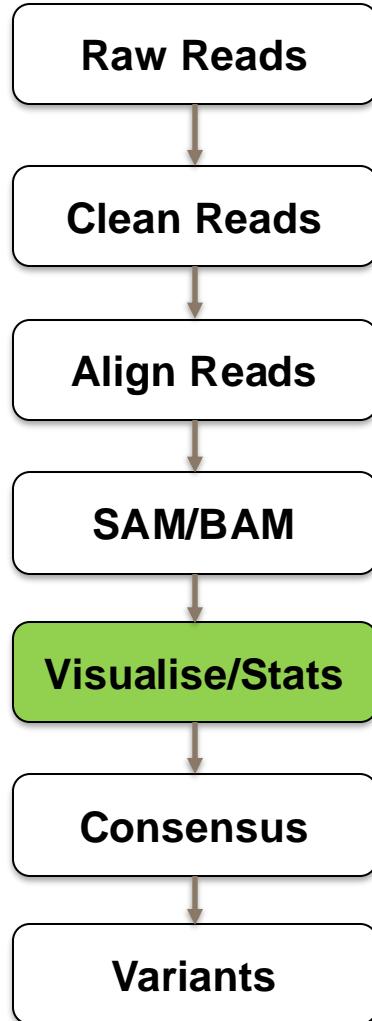
QNAME	FLAG	RNAME	POS	MAPQ	CIGAR	RNEXT	PNEXT	TLEN	SEQ	QUALITY
Read3	10	MyRefSeq	28	52	2M1D4M				ACCAGA	IHGFFF
Read8	4	*	0	0					GCGACAC	IIHHGGG

Host filtering – exploiting flag4 (unmapped)



- Create read files without the human host
- samtools has a host of other function available:
 - samtools fastq
 - **samtools depth**
 - samtools stats
 - samtools ampliconclip
 - **samtools idxstats**
 - **samtools flagstat**
 - samtools consensus

Post Assembly – after the BAM



- Post assembly checks
 - Summary statistics:
 - Number of reads aligned
 - Number of reads unaligned
 - **Average depth of coverage**
 - **Breadth of coverage**
 - **Coverage plot**
 - **Visualisation of entire alignment**

Coverage

	1	2	3
Pos:	1234567890123456789012345678901234		
Ref:	ACGGTGACACGTAGCAGTACGCGGGTTACACAGA		
	ACGG CGA	CAGT TCG	AC- CAGA
	A GACGTA		GC GGGTT
		GTAGCAGT	TTACACAG
	G CGACAC	T CGCGGG	
	CGG CGAC	AGT TCGC	TACACAT T
	ACG -AGC		GGG GTAC
Cov:	122333433333233333433334443331		

Coverage Depth & Breadth

- **Coverage** is the number of reads that “cover” a particular genome coverage
 - **Depth**
- Average (mean) coverage: the average coverage across all genome positions
- Breadth of coverage: how much of the genome is actually covered

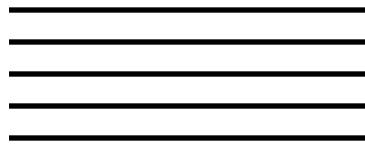
Viral Reference Genome



Average coverage = 1

Breadth = 100%

Viral Reference Genome



Average coverage = 1

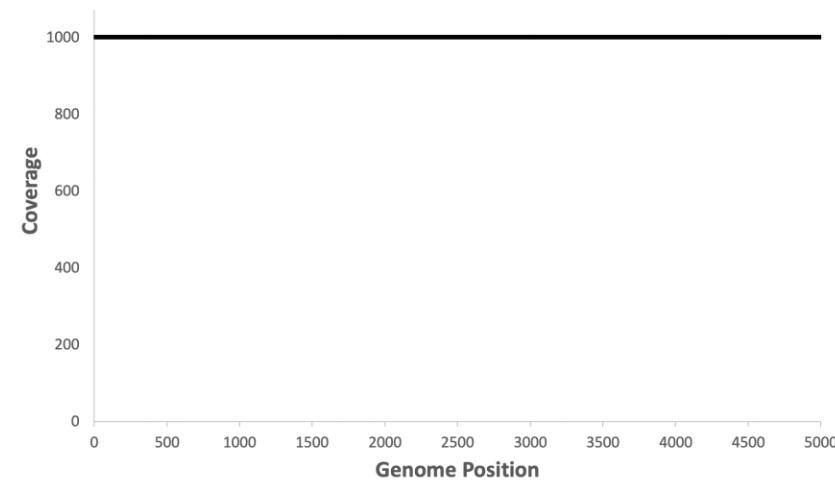
Breadth = 20%

Mode, Median, Quartiles would be different

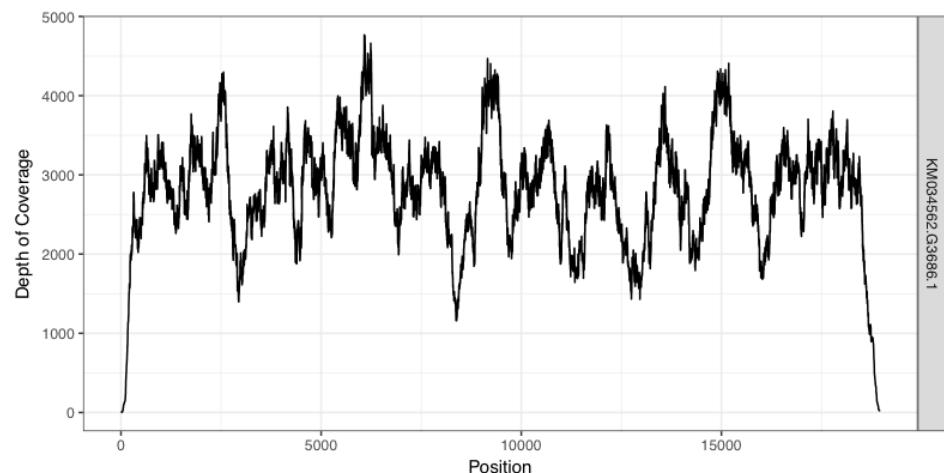
Perfect Coverage Plots

- High uniform coverage across the entire genome
- Biases in library prep fragmentation and PCR (GC content)
 - the terminal ends are typically poorly covered
- Biases in bait capture, amplicon/primer efficiency, extraction methods

Don't think I've ever seen this

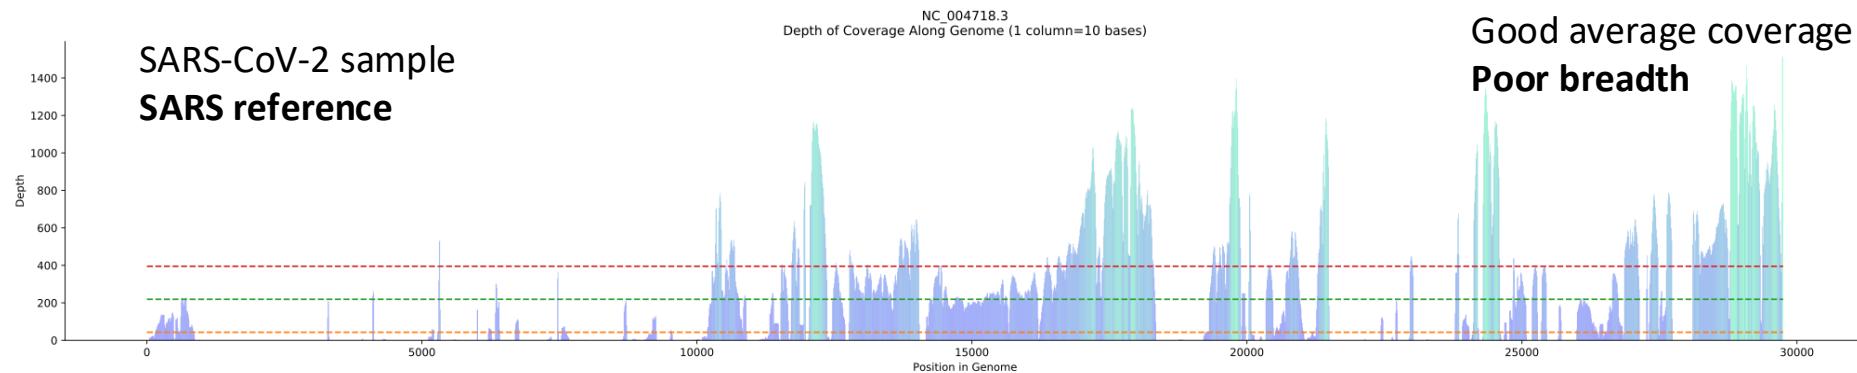


This is perfectly normal

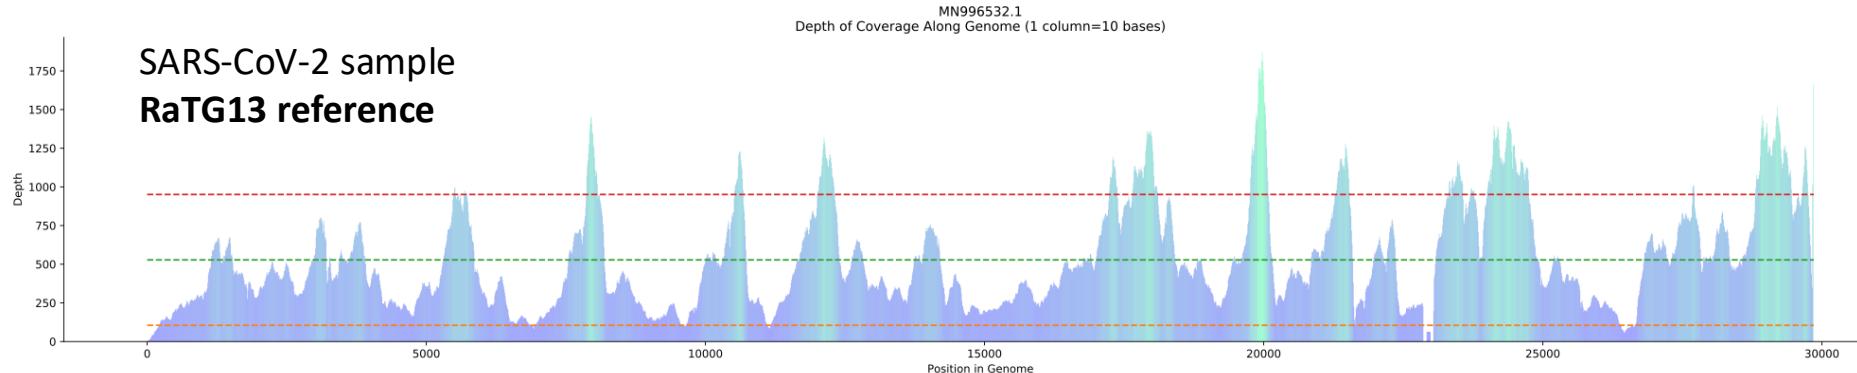


Coverage plots – bad reference

- Sporadic coverage with frequent regions dropping down to zero can indicate a poor reference seq
- The reference is too divergent in many regions and reads can not be aligned at the nucleotide level



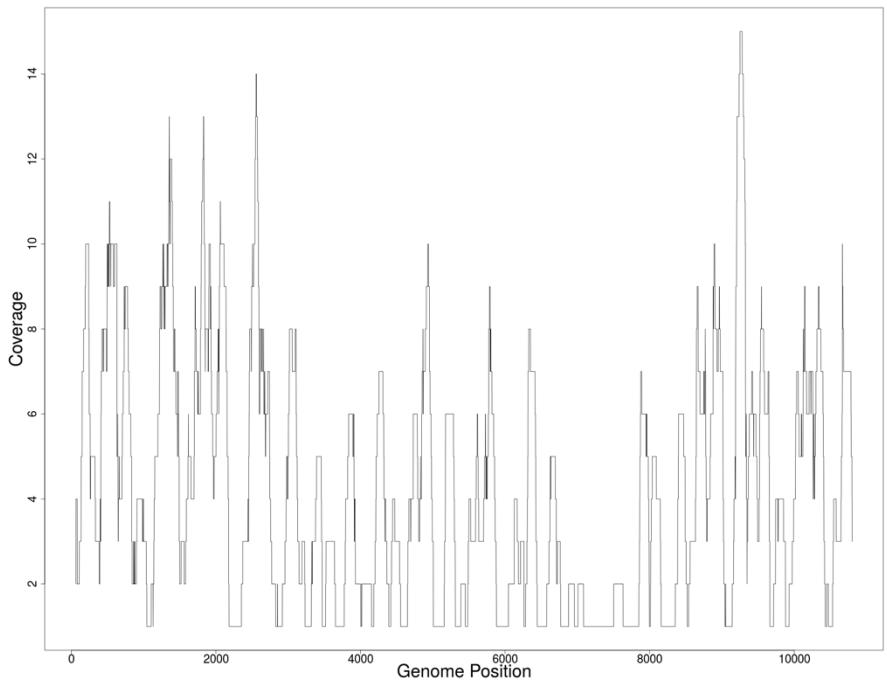
- Viruses can be very diverse – aligning to a different genotype/strain can give obscure results
- Align to different refs, genotype detection tools, **de novo assembly**



Coverage plots – low coverage

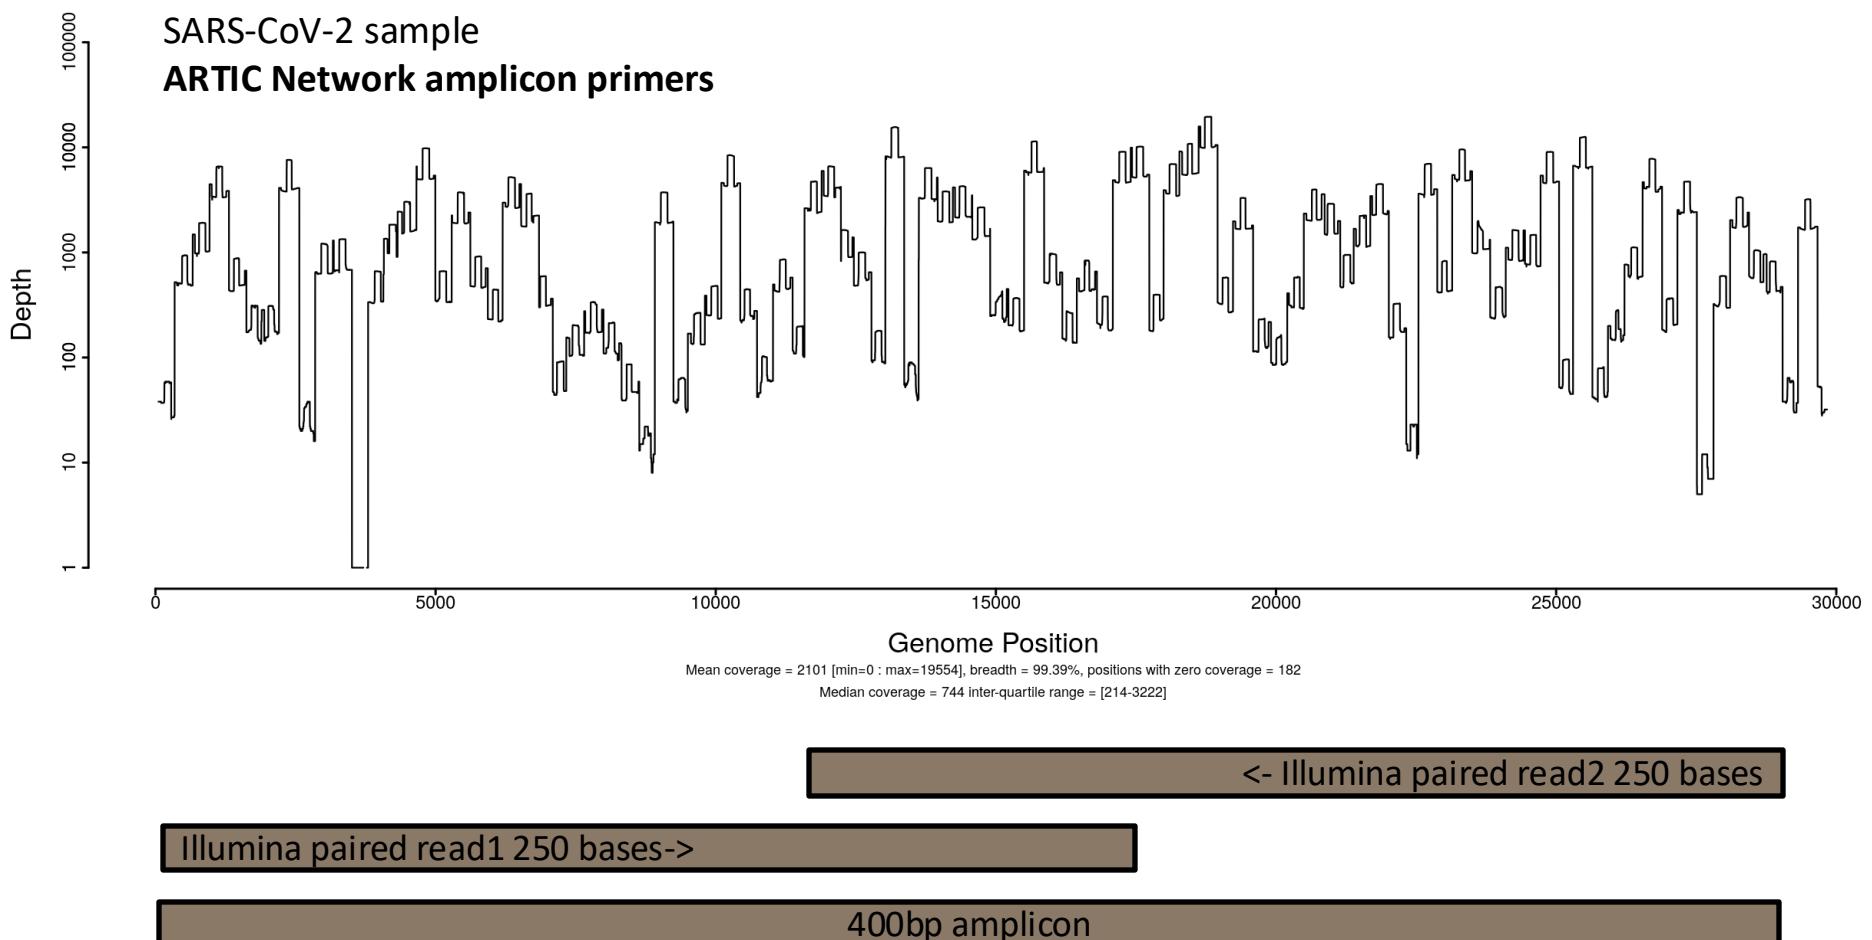
- Sometimes there is just not enough data present
 - Lower read trimming threshold
 - Just use the raw reads
 - Will be noisy
- Re-run the sample
 - Perhaps it was a bad run
 - Combine run data
- PCR amplification
- Bait capture

Louping ill virus sample



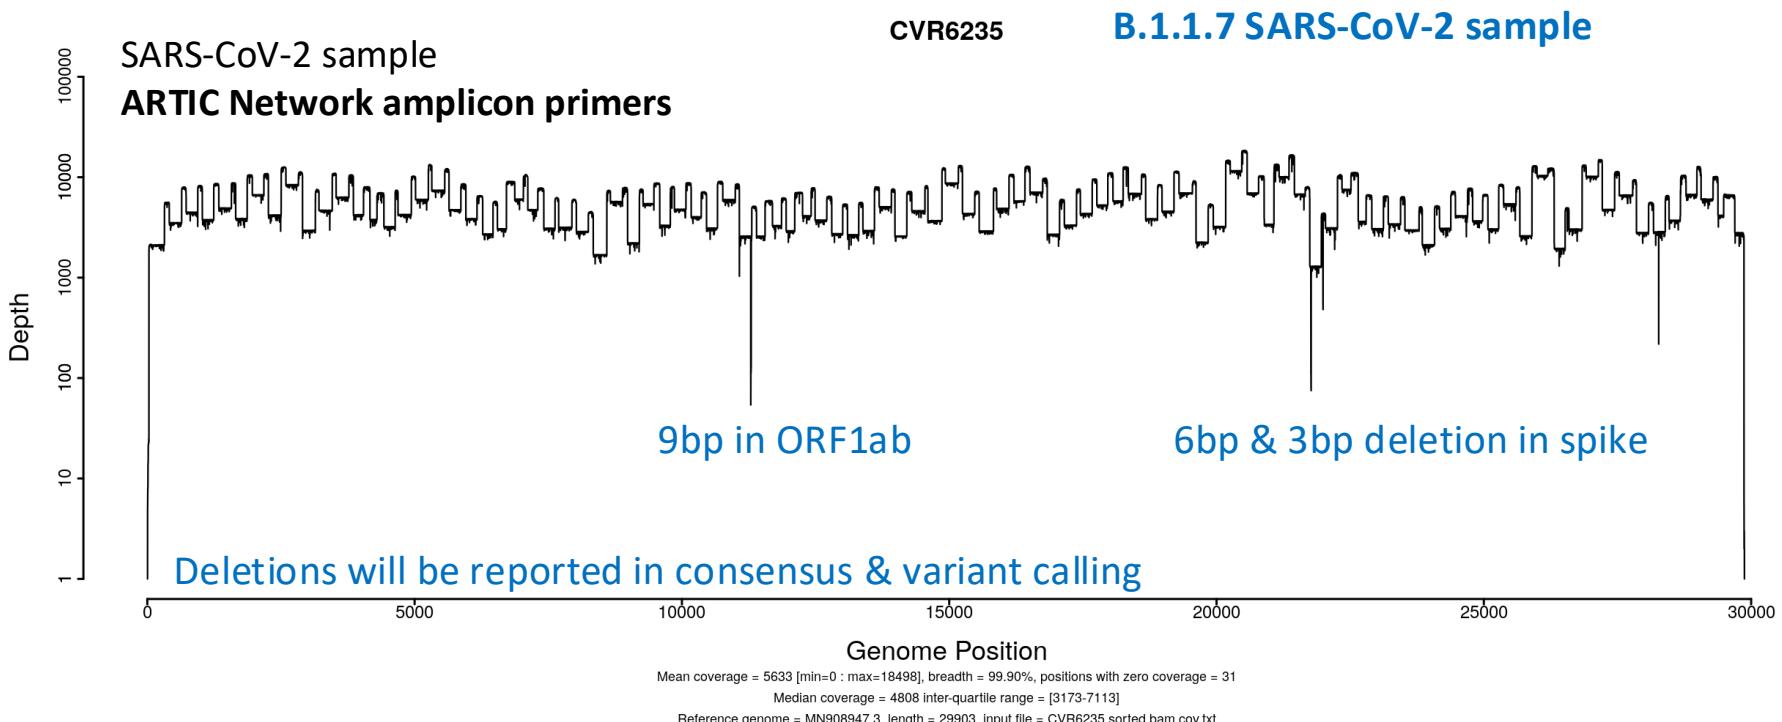
Coverage plots – amplicons

- Amplicon data can give step like plots



Coverage plots – deletions

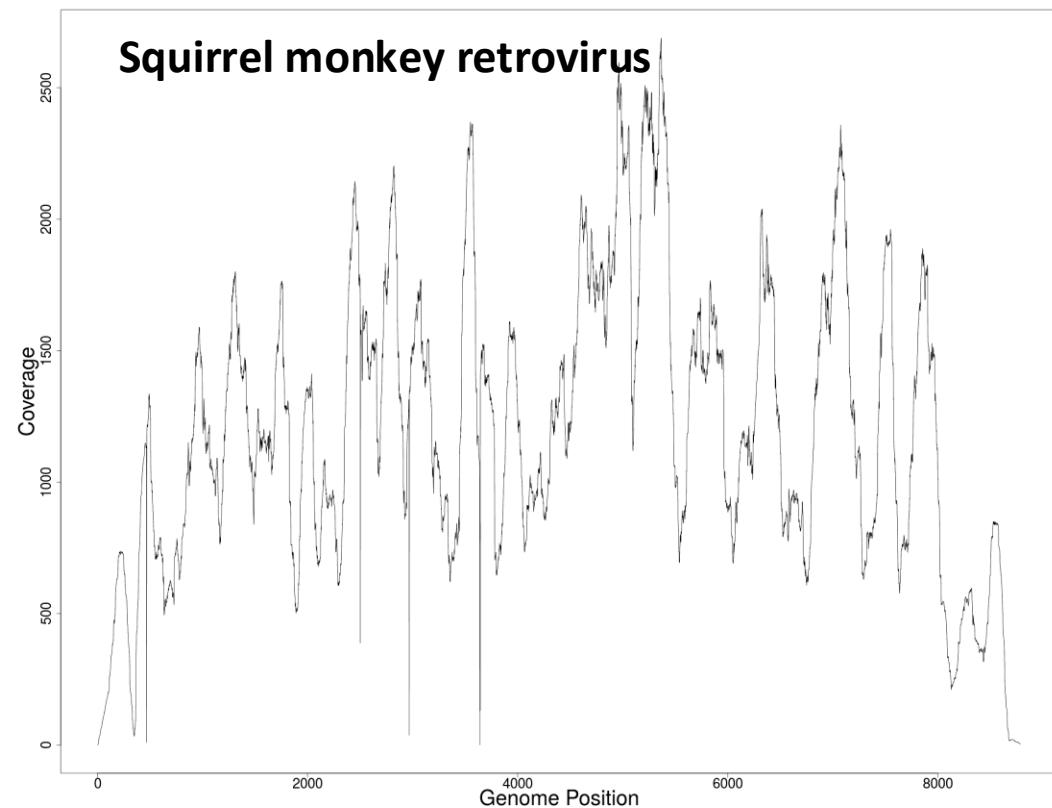
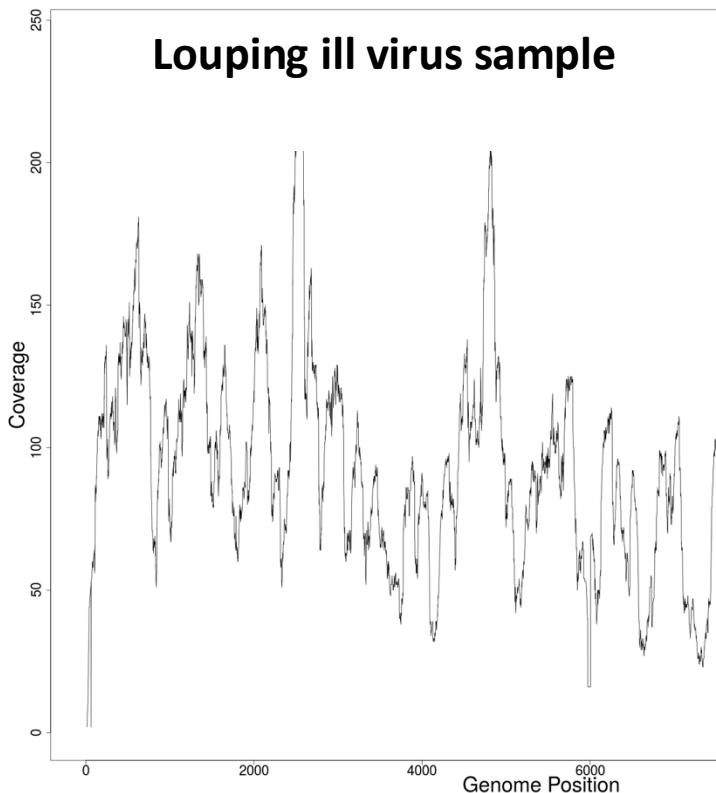
- Sudden drops in coverage at a small number of sites can indicate deletions with respect to the reference



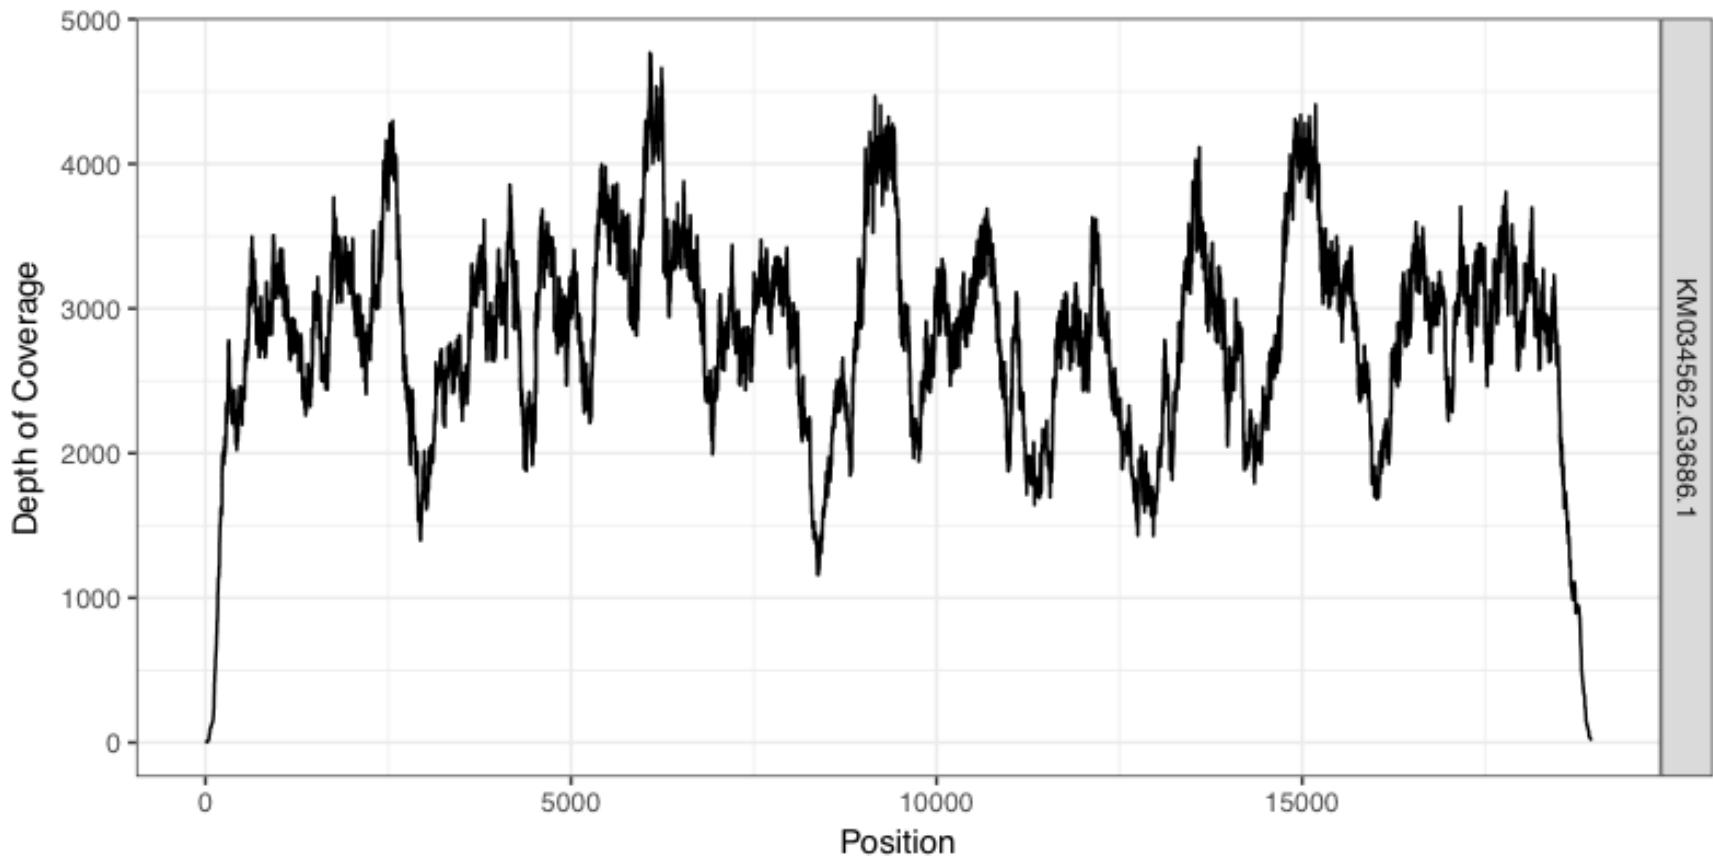
This is a log plot and noisy minion data – the deleted sites still have some coverage but this is nothing compared to the other sites

Reference assembly – tunnel vision

- With reference assembly you automatically focus on a single virus
 - You can align to multiple viruses in one go
 - But you will still need to decide what viruses to investigate
- Good to run kraken/centriguge on your samples to (viral & mycoplasma contaminants)**



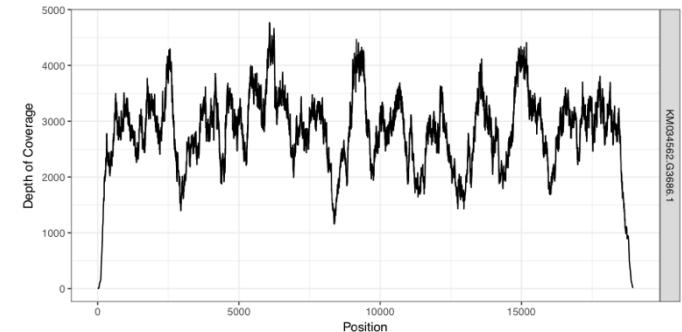
How do you create coverage plots?



samtools depth

- samtools has a built in function called ‘depth’
- `samtools depth -aa -d 0 my.bam > my_depth.txt`
- -aa: output data for absolutely all positions (even positions with zero coverage)
- -d 0: disable the maximum depth to report [default is 8000]
- 3 column text file:

Chromosome	Position	Depth
MN908947.3	1	0
MN908947.3	2	13
MN908947.3	3	34
...

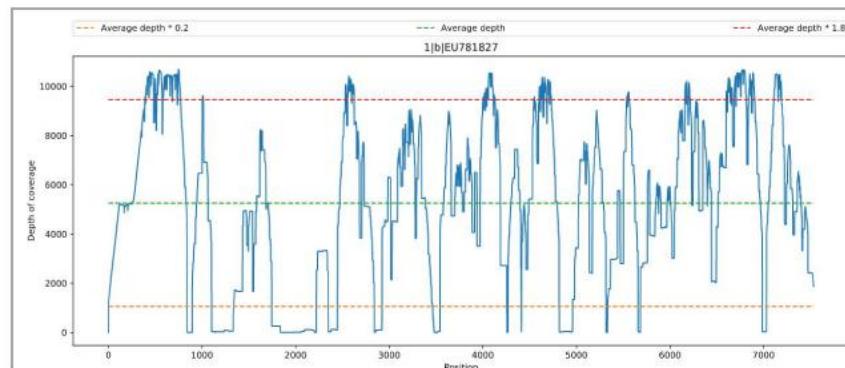


All chromosome will be reported in turn

weeSam - <https://github.com/centre-for-virus-research/weeSAM>

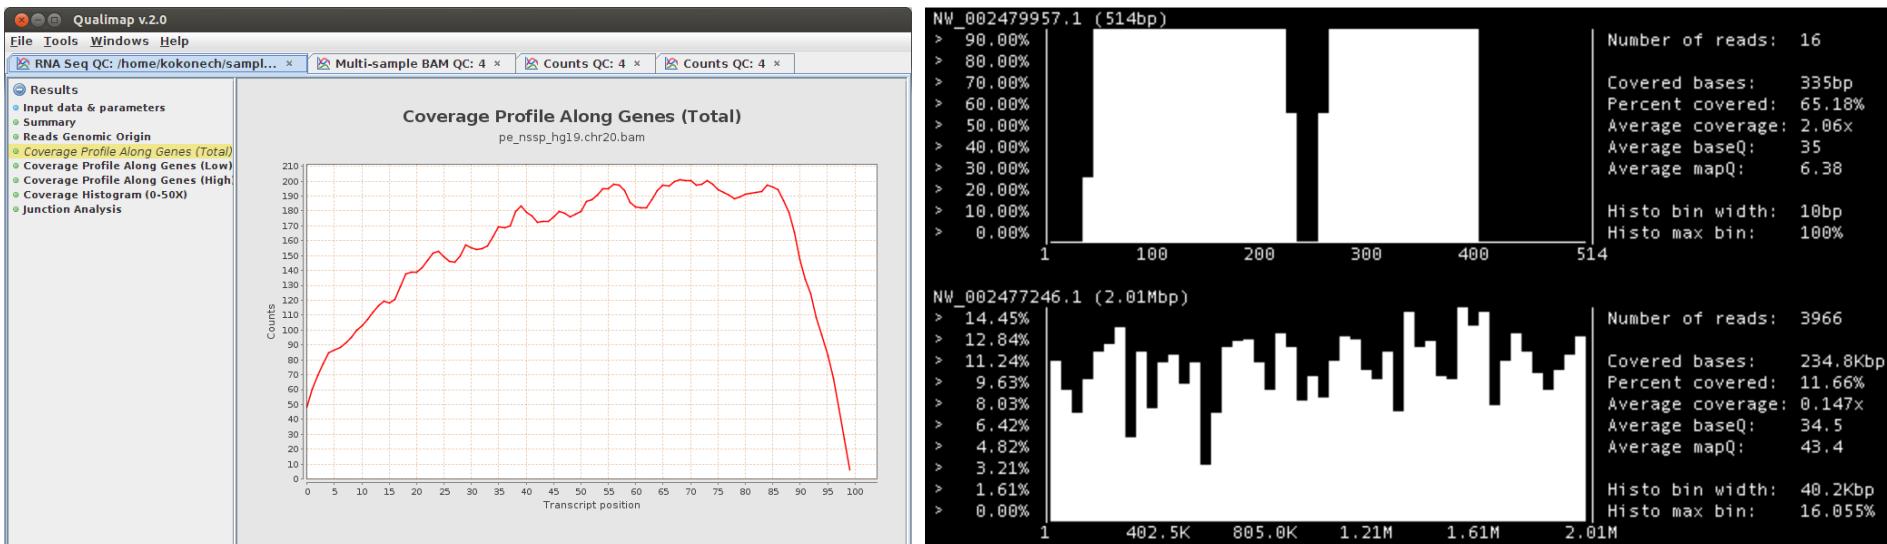
- weeSam is program that can give you information on breadth and depth of coverage as well as generate a coverage plot automatically
- weeSAM --bam 1a.bam --html 1a
- 1a_html_results/1a.html

Ref_Name	Ref_Len	Mapped_Reads	Breadth	%_Covered	Min_Depth	Max_Depth	Avg_Depth	Std_Dev	Above_0.2_Depth	Above_1_Depth	Above_1.8_Depth	Variation_Coefficient
NC_004102.1 Hepatitis C virus genotype 1, complete genome	9646	640000	9646	100.00	13	10729	9941.89	1699.34	98.82	90.91	0.00	0.17

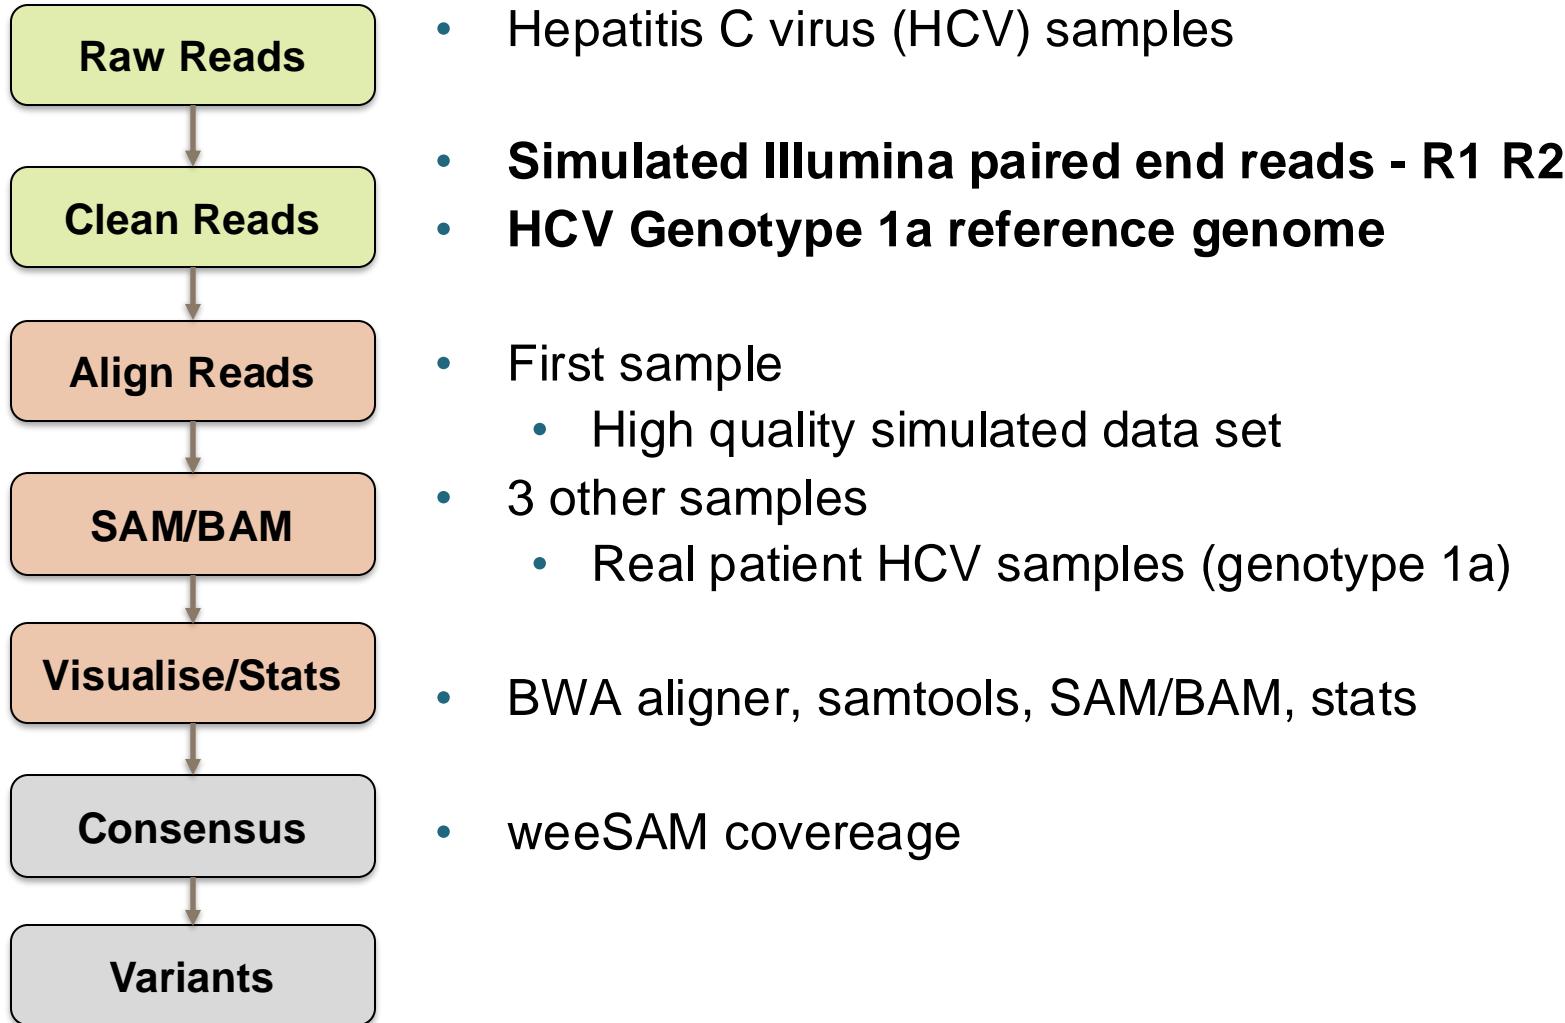


Other tools

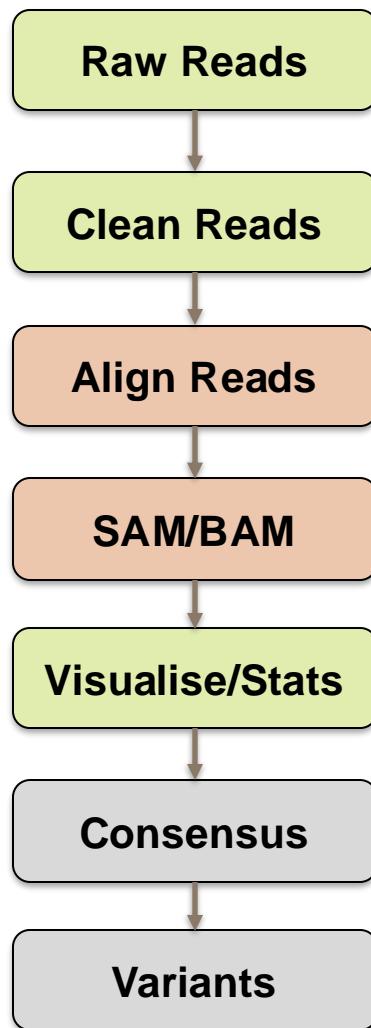
- Qualimap: <http://qualimap.conesalab.org>
- bamCov – <https://github.com/fbreitwieser/bamcov>



Practical



Practical – HCV_SIM commands – adapt for another sample



```
bwa index ../1a_hcv_ref.fasta
```

```
bwa mem ../1a_hcv_ref.fasta hcv_sim_R1.fq  
hcv_sim_R2.fq > 1b.sam
```

```
samtools view -b 1b.sam
```

```
samtools sort -o 1b.bam
```

```
samtools index 1b.bam
```

```
rm 1b.sam
```

```
samtools view -c -F4 1b.bam
```

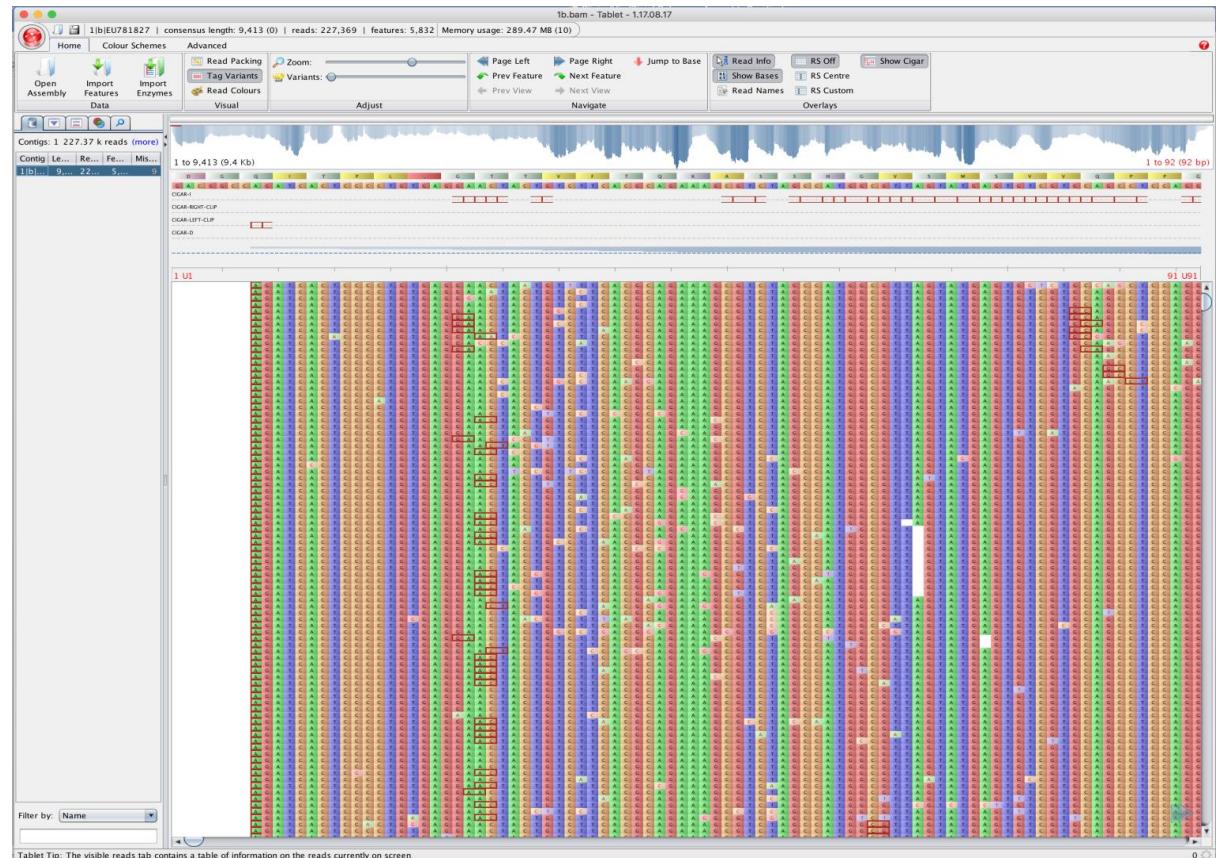
```
samtools view -c -f4 1b.bam
```

The End ... Tablet

- Tablet demo if time later on

Tablet: <https://ics.hutton.ac.uk/tablet/>

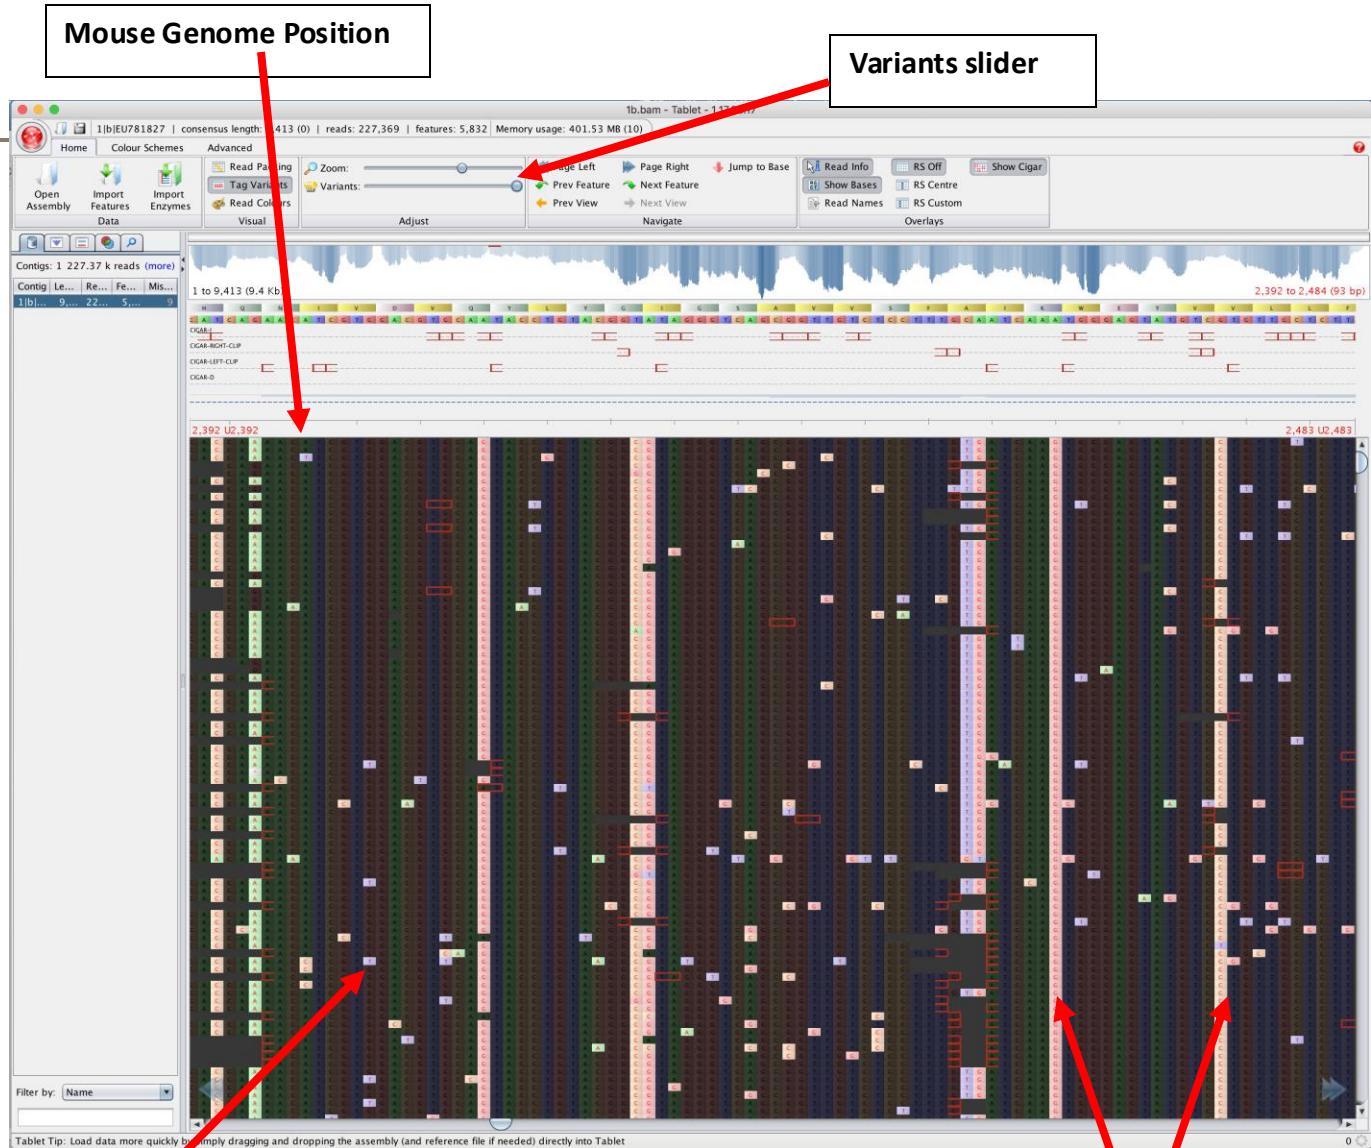
- **tablet**
- Zoom, scroll, colour schemes: nucleotides, direction, mutations
- Tablet is a lightweight, high-performance graphical viewer for next generation sequence assemblies and alignments.
 - BAM file
 - Reference file



Tablet



Tablet



**Minority
variants**

**Consensus level
variants**